

Workflow reliability - incomplete delivery

Branch: shbs-microsoft-workflow-app-verification
Current scope: sample-app verification and prioritized workflow reliability fixes.
[Design](#) · [Continuation plan](#) · [Acceptance ledger](#)

Outcome

Partial delivery: reliability, five-verb CLI, tool bindings, README/guides and the native parser core are committed. Native serialization/candidates and two corrected reference/dependency defects now have [bounded preservation approval](#). A recurring grouped-test timeout remains unresolved. Collectors, guarded metrics, recovery and release quality are incomplete.

Implemented and pending

Delivered: source-bound atomic reports, exact argv/cwd, assertion-only kills, freshness/restoration and fail-closed completeness; redesigned README/banner, user guide and installation/contributor/architecture updates. Docs were rendered locally at desktop/mobile sizes in light/dark modes.

C3 integration: current-design/adoption, earlier compatibility, research supersession and authoritative status guidance; README navigation; three CLI regressions in ordinary discovery. [Proof](#) and [review](#) retain runtime limits.

Live-app distinction: the existing planner now passes a fresh native/browser regression. A new agent-driven app build exercising the updated workflow has not run. Neither app regression nor parser tests prove all 29 workflow features. Auxiliary C4 preparation remains on hold.

Acceptance and gates

AC	Current verdict	Evidence or blocker
01 - ESM/CJS	VERIFIED-WITH-LIMITATIONS	Windows/Node24 native discovery/routing.
02 - truthful metrics	VERIFIED-WITH-LIMITATIONS	Final correction approved ; complete quality still due.
03-04 - guards	Ordinary CLI verified with limits	Real admission/nonmetric receipts/adoption; no metric closure or recovery.
05-08 - resume/contracts/research/status	Partial	Ordinary behavior and C3 guidance verified with limits; live-host work remains.
09 - delivery	BLOCKED	Continuation gates and actual measurements due.
10 - documentation	Verified with limitations	README/guides, links and local browser proof independently accepted.

G1	G2	G3	G4	G5	G6
Pass	Autonomous acceptance; human unconfirmed	Continuation accepted	Corrected cases pass; grouped timeout	Bounded proof only; fresh-agent workflow pending	Partial preservation; quality blocked

Measurements, not substitutes

Observation	Result	Scope
Historical foundation tests	I1: 54; C1/Q1: 121; C2b-S: 110 passed	Separate I1 , C1/Q1 and state-seam runs; not additive or current-tree proof.
Joint checkpoint before R1	262/263 Python; native13 passed	One symlink case environment-unverified ; native includes five Python replays.
R1 fresh affected proof	10 methods passed	20 malformed manifests reject; valid control runs . Earlier262 not refreshed or added.
C3 independent proof	CLI33; native70; 12 browser views	217 inputs unchanged ; five native entries replay CLI cases. Three navigation actions; 20 installed-layout links. Not added to earlier counts.
Historical parser-core proof	15 methods + separate two passed	3,593 inputs unchanged ; earlier bytes, not current full-suite or pipeline proof.
Serialization correction proof	Seven distinct methods have passing evidence over eight invocations	223 source/history pins unchanged . Grouped run failed on recurring candidate idle timeout; sole unchanged-byte retry passed. Reliable combined execution remains unproven.
Fresh existing-planner regression	37 tests; four flows/126 checks passed	99.33% lines / 96.92% branches / 96.63% functions ; browser replay . Unchanged app inputs; one native capture correction/rerun, not a product fix.
Historical evaluation checks	Native65; corpus10/79 assertions; CSV9/9 graded kills	Native / fixed corpus ; seven mutants ungraded, 0.9 bar unchanged. Not live-model or whole-change proof.
Actual quality	No current complete proof	Historical incomplete report ; subsequent edits invalidate current-snapshot use.

Further [measurement implementation](#) is paused for the app-driven exercise. Refusal tests, installed tools and dry runs cannot replace actual measurements.

Quality scorecard

Dimension	Status	Basis
Correctness	Bounded CLI/bindings accepted	Independent proof/review; environment and full-quality limits retained.
Grounding	Exhaustive audit unverified	Cited research and source-backed review.
Design fidelity	Bounded changes reviewed	Approved partial contracts honored; remaining requirements not waived.
Scope fidelity	Historical I1: 4/5 bounded review	Continuation not rescored; overall request incomplete.
Reuse / DRY	Unverified	Duplication proof unavailable.
Modularity	Unverified	Complexity/cycles/architecture unavailable.
Maintainability	Not numerically rescored	Command-selection duplication resolved.
Robustness	Unverified	Static proof unavailable.
Test quality	Unverified for ship	Whole-change mutation/coverage missing.

[I1 re-review](#) approves limited preservation, not release. Continuation scores remain unproven.

How to finish

Preserve this bounded serialization checkpoint, then prioritize actual app-driven workflow verification rather than another internal expansion. Earlier parser/app evidence is in `886d699` and `8881a7e` . Remaining command/manifests/replay, collectors, guarded binding and Q3-Q5 are not waived. Auxiliary C4 preparation stays held; recovery and publication remain blocked.

Human approval remains unconfirmed. Executed runtimes: Python 3.14.2 / Node 24.11.1 on Windows; no broader matrix proof. Hash scope is not a sandbox; producer identities are self-declared. Compaction is not fresh-agent resume. Publication remains blocked by EMU 403; no retry or bypass occurred.